

Mini Project 1

Problem Statement: Which model is suitable for Insurance dataset

Importing Packages

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 from sklearn.model_selection import train_test_split
        6 from sklearn.linear_model import LinearRegression, LogisticRegression
        7 from sklearn.metrics import r2_score
```

Read data

```
In [2]: 1 df=pd.read_csv(r"C:\Users\P. VIJAY KUMAR\Downloads\insurance.csv")
        2 df
```

```
Out[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

Data Pre-processing

```
In [3]: 1 df.isnull().any()
```

```
Out[3]: age          False
sex          False
bmi          False
children     False
smoker       False
region       False
charges      False
dtype: bool
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype
='object')
```

```
In [5]: 1 df["sex"].value_counts()
```

```
Out[5]: sex
male      676
female    662
Name: count, dtype: int64
```

```
In [6]: 1 convert={"sex":{"male":1,"female":2}}
2 df=df.replace(convert)
3 df
```

```
Out[6]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520
...
1333	50	1	30.970	3	no	northwest	10600.54830
1334	18	2	31.920	0	no	northeast	2205.98080
1335	18	2	36.850	0	no	southeast	1629.83350
1336	21	2	25.800	0	no	southwest	2007.94500
1337	61	2	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [7]: 1 df["smoker"].value_counts()
```

```
Out[7]: smoker
no      1064
yes      274
Name: count, dtype: int64
```

```
In [8]: 1 convert={"smoker":{"no":1,"yes":2}}
2 df=df.replace(convert)
3 df
```

```
Out[8]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	2	southwest	16884.92400
1	18	1	33.770	1	1	southeast	1725.55230
2	28	1	33.000	3	1	southeast	4449.46200
3	33	1	22.705	0	1	northwest	21984.47061
4	32	1	28.880	0	1	northwest	3866.85520
...
1333	50	1	30.970	3	1	northwest	10600.54830
1334	18	2	31.920	0	1	northeast	2205.98080
1335	18	2	36.850	0	1	southeast	1629.83350
1336	21	2	25.800	0	1	southwest	2007.94500
1337	61	2	29.070	0	2	northwest	29141.36030

1338 rows × 7 columns

```
In [9]: 1 df["region"].value_counts()
```

```
Out[9]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [10]:

```
1 convert={"region":{"southeast":1,"southwest":2,"northwest":3,"northeast":4}
2 df=df.replace(convert)
3 df
```

Out[10]:

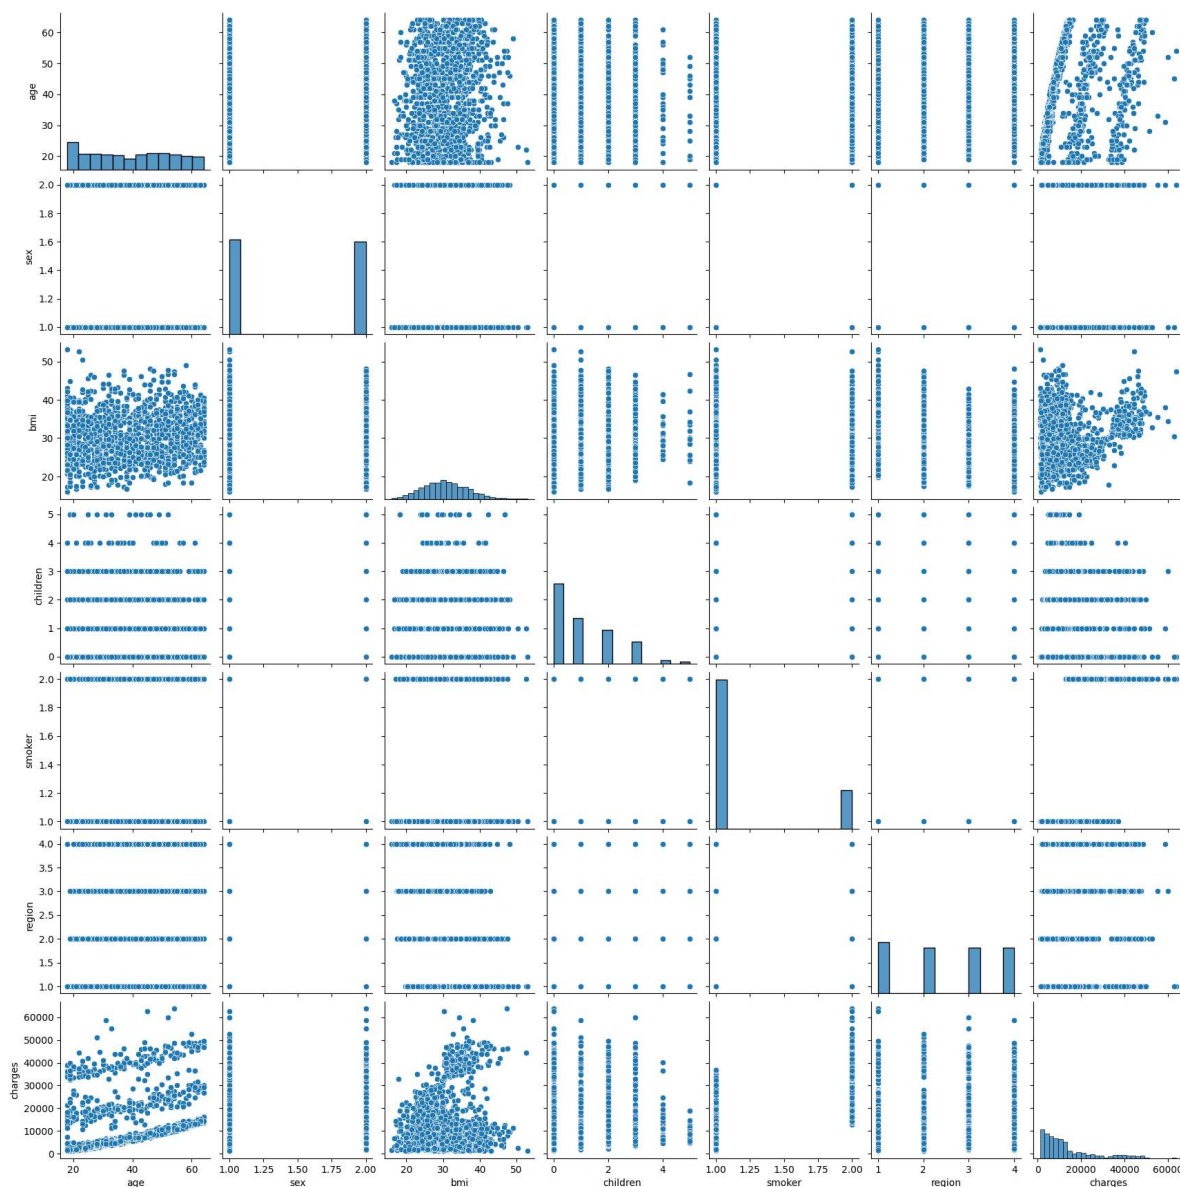
	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	2	2	16884.92400
1	18	1	33.770	1	1	1	1725.55230
2	28	1	33.000	3	1	1	4449.46200
3	33	1	22.705	0	1	3	21984.47061
4	32	1	28.880	0	1	3	3866.85520
...
1333	50	1	30.970	3	1	3	10600.54830
1334	18	2	31.920	0	1	4	2205.98080
1335	18	2	36.850	0	1	1	1629.83350
1336	21	2	25.800	0	1	2	2007.94500
1337	61	2	29.070	0	2	3	29141.36030

1338 rows × 7 columns

Data Visualization


```
In [11]: 1 sns.pairplot(df)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x25ddca779a0>
```




```
In [12]: 1 columns_to_plot=df[['age', 'sex', 'bmi', 'children', 'smoker', 'region',
2   sns.pairplot(columns_to_plot,x_vars=['age', 'sex', 'bmi', 'children', 'region',
3   aspect=1,kind='scatter')]
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x25d837402e0>
```



```
In [13]: 1 #To plot heatmap to find out correlations
2 subset_data =df[['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges']]
3 correlation_matrix = subset_data.corr() # Use .corr() for correlation
4 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
```

Out[13]: <Axes: >



Feature Scaling: Splitting data into training and testing

```
In [14]: 1 X=np.array(df["smoker"]).reshape(-1,1)
2 y=np.array(df["charges"]).reshape(-1,1)
```

```
In [15]: 1 x_train,x_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_
```

Applying Linear Regression

```
In [16]: 1 lr=LinearRegression()
```

```
In [17]: 1 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [18]: 1 lr.score(x_test,y_test)
```

Out[18]: 0.6197902354385714

```
In [19]: 1 sns.lmplot(x="charges",y="smoker",data=df,order=2,ci=None)
```

Out[19]: <seaborn.axisgrid.FacetGrid at 0x25d834081f0>



Applying Logistic Regression

```
In [20]: 1 x=np.array(df["charges"]).reshape(-1,1)
          2 y=np.array(df["smoker"]).reshape(-1,1)
          3 df.dropna(inplace=True)
```

```
In [21]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_s
```

```
In [22]: 1 lg=LogisticRegression(max_iter=10000)
```

```
In [23]: 1 lg.fit(x_train,y_train)
```

C:\Users\P. VIJAY KUMAR\AppData\Roaming\Python\Python310\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

Out[23]: LogisticRegression(max_iter=10000)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [24]: 1 lg.score(x_test,y_test)
```


Out[24]: 0.8930348258706468

```
In [25]: 1 prediction=lg.predict(x_test)
          2 r2=r2_score(y_test,prediction)
          3 r2
```

Out[25]: 0.335179416176301


```
In [28]: 1 sns.regplot(x=x,y=y,data=df,logistic=True,ci=None)
```

```
Out[28]: <Axes: >
```



Applying Decision Tree

```
In [29]: 1 from sklearn.tree import DecisionTreeClassifier
2 dtl=DecisionTreeClassifier(random_state=0)
```

```
In [30]: 1 dtl.fit(x_train,y_train)
```

```
Out[30]: DecisionTreeClassifier(random_state=0)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [32]: 1 dtl.score(x_test,y_test)
```

```
Out[32]: 0.8880597014925373
```

Applying Random Forest

```
In [33]: 1 from sklearn.ensemble import RandomForestClassifier
        2 rfc=RandomForestClassifier()
```

```
In [34]: 1 rfc.fit(x_train,y_train)
```

C:\Users\P. VIJAY KUMAR\AppData\Local\Temp\ipykernel_17540\4070307935.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
rfc.fit(x_train,y_train)

Out[34]: RandomForestClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [35]: 1 rfc.score(x_test,y_test)
```

Out[35]: 0.8880597014925373

CONCLUSION: Here i developed LinearRegression model, LogisticRegression model, Decision Tree model and RandomForest model for provided dataset.Among them the LogisticRegression has got more accuracy on given dataset, So LogisticRegression is best fittedmodel for our dataset.

```
In [ ]: 1
```